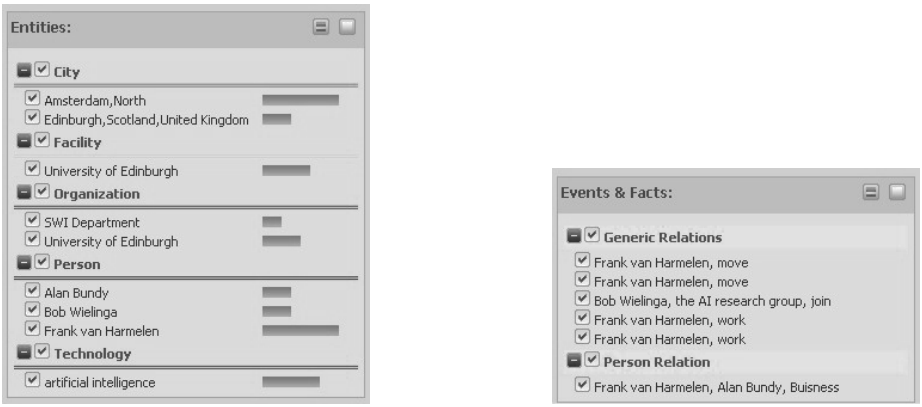

Figure 6.5: Objects and relations recognized by OpenCalais

tured in natural language texts. It does this by analyzing a text, recognizing named entities such as people, locations, companies, etc., and annotating the text with RDF to identify these named entities.

As a simple example, the following text produces the entities and relations as shown in figure 6.5:

> After studying mathematics and computer science in Amsterdam, Frank van Harmelen moved to the Department of AI of the University of Edinburgh. While in Edinburgh, he worked with Prof. Alan Bundy on proof planning for inductive theorem proving. After his PhD research, he moved back to Amsterdam where he worked from 1990 to 1995 in the SWI Department under Prof. Bob Wielinga. In 1995 he joined the AI research group at the Vrije Universiteit Amsterdam.

For these objects, Reuters uses its own URIs to identify the entities, but many of these private URIs are linked through owl:sameAs statements to more widely adopted URIs for the same entities taken from GeoNames, DBPedia, FreeBase, the Internet Movie Database, and others. This allows us to further retrieve information about these entities from Reuter's database, such as financial information, names of board mem-